

```
import os
os.makedirs("/root/.kaggle", exist_ok=True)

from google.colab import files
files.upload()

!cp /content/kaggle.json /root/.kaggle/
!chmod 600 /root/.kaggle/kaggle.json
```

kaggle.json
kaggle.json(application/json) - 72 bytes, last modified: 23/09/2025 - 100% done
 Saving kaggle.json to kaggle.json

```
!kaggle datasets download -d mlg-ulb/creditcardfraud
!unzip creditcardfraud.zip
```

Dataset URL: <https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>
 License(s): DbCL-1.0
 Downloading creditcardfraud.zip to /content
 0% 0.00/66.0M [00:00<?, ?B/s]
 100% 66.0M/66.0M [00:00<00:00, 1.39GB/s]
 Archive: creditcardfraud.zip
 inflating: creditcard.csv

Project Explanation – Credit Card Fraud Detection (Data Exploration)

In this project, I started by performing an exploratory data analysis (EDA) on the Credit Card Fraud Detection dataset, which contains anonymized transaction data. The dataset includes numerical features (V1–V28), the transaction amount, the transaction time, and a target column Class (0 = legitimate transaction, 1 = fraudulent transaction).

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score, roc_curve
from imblearn.over_sampling import SMOTE
```

I used the Pandas library to load the dataset into a DataFrame for further analysis. Displayed the first five rows of the dataset to understand the structure and check column names.

Verified that the target column is labeled as Class. The dataset contains 284,807 rows and 31 columns.

This shows that we are working with a large-scale dataset, which is important for building reliable machine learning models.

Provided a statistical overview of the numerical columns (mean, min, max, standard deviation, etc.).

This helped to identify the ranges and distributions of different features.

```
import pandas as pd

csv_path = "/content/creditcard.csv"
df = pd.read_csv(csv_path)

print(df.head())

Time      V1      V2      V3      V4      V5      V6      V7  \
0  0.0 -1.359807 -0.072781  2.536347  1.378155 -0.338321  0.462388  0.239599
1  0.0  1.191857  0.266151  0.166480  0.448154  0.060018 -0.082361 -0.078803
2  1.0 -1.358354 -1.340163  1.773209  0.379780 -0.503198  1.800499  0.791461
3  1.0 -0.966272 -0.185226  1.792993 -0.863291 -0.010309  1.247203  0.237609
4  2.0 -1.158233  0.877737  1.548718  0.403034 -0.407193  0.095921  0.592941

          V8      V9      ...      V21      V22      V23      V24      V25  \
0  0.098698  0.363787  ... -0.018307  0.277838 -0.110474  0.066928  0.128539
1  0.085102  -0.255425  ... -0.225775 -0.638672  0.101288 -0.339846  0.167170
2  0.247676 -1.514654  ...  0.247998  0.771679  0.909412 -0.689281 -0.327642
3  0.377436 -1.387024  ... -0.108300  0.005274 -0.190321 -1.175575  0.647376
```

```
4 -0.270533 0.817739 ... -0.009431 0.798278 -0.137458 0.141267 -0.206010
```

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

[5 rows x 31 columns]

```
!ls /content
```

```
creditcard.csv creditcardfraud.zip kaggle.json sample_data
```

```
import pandas as pd

df = pd.read_csv("/content/creditcard.csv")

print(df.head())

print("Shape:", df.shape)

print(df.describe())

print(df['Class'].value_counts())
```

[5 rows x 31 columns]

Shape: (284807, 31)

	Time	V1	V2	V3	V4	\
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	
	V5	V6	V7	V8	V9	\
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	
mean	9.604066e-16	1.487313e-15	-5.556467e-16	1.213481e-16	-2.406331e-15	
std	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00	
min	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+01	
25%	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01	
50%	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02	
75%	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01	
max	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+01	
	V21	V22	V23	V24	\	
count	...	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	
mean	...	1.654067e-16	-3.568593e-16	2.578648e-16	4.473266e-15	
std	...	7.345240e-01	7.257016e-01	6.244603e-01	6.056471e-01	
min	...	-3.483038e+01	-1.093314e+01	-4.480774e+01	-2.836627e+00	
25%	...	-2.283949e-01	-5.423504e-01	-1.618463e-01	-3.545861e-01	
50%	...	-2.945017e-02	6.781943e-03	-1.119293e-02	4.097606e-02	
75%	...	1.863772e-01	5.285536e-01	1.476421e-01	4.395266e-01	
max	...	2.720284e+01	1.050309e+01	2.252841e+01	4.584549e+00	

```
[8 rows x 31 columns]
Class
0    284315
1     492
Name: count, dtype: int64
```

```
print("Shape of dataset:", df.shape)
print(df['Class'].value_counts())
```

```
Shape of dataset: (284807, 31)
Class
0    284315
1     492
Name: count, dtype: int64
```

```
X = df.drop("Class", axis=1)
y = df["Class"]
```

The dataset is highly imbalanced:

Most transactions are legitimate (Class 0).

A very small fraction are fraudulent (Class 1).

This imbalance is a critical challenge in fraud detection and requires special handling (e.g., resampling, anomaly detection, or cost-sensitive learning).

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42, stratify=y)
```

```
print("Train size:", X_train.shape, " Test size:", X_test.shape)
```

```
Train size: (227845, 30) Test size: (56962, 30)
```

```
sm = SMOTE(random_state=42)
X_res, y_res = sm.fit_resample(X_train, y_train)
```

```
print("After SMOTE:", y_res.value_counts())
```

```
After SMOTE: Class
0    227451
1    227451
Name: count, dtype: int64
```

```
log_reg = LogisticRegression(max_iter=1000)
log_reg.fit(X_res, y_res)
y_pred_lr = log_reg.predict(X_test)
```

```
print("---- Logistic Regression ----")
print(classification_report(y_test, y_pred_lr))
print("ROC-AUC:", roc_auc_score(y_test, log_reg.predict_proba(X_test)[:,1]))
```

	precision	recall	f1-score	support
0	1.00	0.97	0.99	56864
1	0.06	0.92	0.11	98
accuracy			0.97	56962
macro avg	0.53	0.95	0.55	56962
weighted avg	1.00	0.97	0.99	56962

```
ROC-AUC: 0.970987885165321
```

```
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_res, y_res)
y_pred_rf = rf.predict(X_test)
```


```
print("---- Random Forest ----")
print(classification_report(y_test, y_pred_rf))
print("ROC-AUC:", roc_auc_score(y_test, rf.predict_proba(X_test)[:,1]))
```

```
fig, ax = plt.subplots(1, 2, figsize=(12,5))

sns.heatmap(confusion_matrix(y_test, y_pred_lr), annot=True, fmt="d", cmap="Blues", ax=ax[0])
ax[0].set_title("Logistic Regression")


sns.heatmap(confusion_matrix(y_test, y_pred_rf), annot=True, fmt="d", cmap="Greens", ax=ax[1])
ax[1].set_title("Random Forest")

plt.show()
```



```
fpr1, tpr1, _ = roc_curve(y_test, log_reg.predict_proba(X_test)[:,1])
fpr2, tpr2, _ = roc_curve(y_test, rf.predict_proba(X_test)[:,1])

plt.figure(figsize=(6,6))
plt.plot(fpr1, tpr1, label="Logistic Regression")
plt.plot(fpr2, tpr2, label="Random Forest")
plt.plot([0,1],[0,1],"--", color="gray")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve")
plt.legend()
plt.show()
```



Double-click (or enter) to edit

T B I <> ⌂ 99 ≡ ≡ — ψ ☺ ☻

Conclusion:

At this stage, I successfully:

Loaded and inspected the dataset.

Understood the feature set and target variable.

Identified class imbalance as a key issue.

Conclusion:

At this stage, I successfully:

Loaded and inspected the dataset.

Understood the feature set and target variable.

Identified class imbalance as a key issue.